

# TAE Week 3 Coaching Guide


LLMs, Attention & Tiny Transformers Capstone

The AI Engineer Program

November 15, 2025

## Abstract

Week 3 connects transformer theory with practical LLM engineering. You will deepen your understanding of attention and tiny decoder-only transformers using the dedicated handout, then embed these ideas into an engineering workflow using Part II of the AI/ML Engineering book and the LLM guide. The capstone centers on a from-scratch tiny transformer language model, built and trained in a Colab notebook, with clear logging, checkpoints, and sampling utilities.



## Contents

<b>1 Purpose &amp; Scope (Week 3)</b>	1
<b>2 Core Resources</b>	1
<b>3 What You Will Learn</b>	3
<b>4 How to Study</b>	3
<b>5 Focus First (Priorities)</b>	4
<b>6 Key Concepts — Concise Recap</b>	4
<b>7 Review Questions</b>	5
<b>8 Practice Checklist</b>	5
<b>9 Tiny Study Loop</b>	6
<b>10 Daily Cadence (Example)</b>	6
<b>11 Capstone Project (Week 3)</b>	6

## 1 Purpose & Scope (Week 3)

This guide coaches you through Week 3. It focuses on the **LLM book and deck**, the **Engineering companion** (Part II), and the **Attention & Tiny Transformers handout** that defines the capstone: a small decoder-only transformer language model implemented from scratch.

### Week 3 Outcomes

By the end of the week you should be able to: (1) explain and implement scaled dot-product attention and (multi-head) self-attention, (2) assemble a tiny decoder-only transformer with positional encodings and causal masking, (3) train and sample from a miniature language model in Colab, and (4) wrap this work in a simple but honest engineering workflow (tracking, checkpoints, and reproducible runs).

## 2 Core Resources

Keep these open while you work; move between them rather than reading any single one linearly.

- **LLM Book (llm.html)** [1] — mindset, architecture, and training pipeline for building a GPT-style model from scratch.
- **LLM Slides (llm\_slides.pdf)** [3] — high-level structure and “why” for each component; skim before and after working in the handout.
- **AI/ML Engineering (eng.html, Part II)** [2] — tracking, configuration, pipelines, and ML lifecycle practices you will apply to the capstone.
- **Engineering Slides (eng\_slides.pdf, Part II)** [4] — condensed view of ML engineering in practice; use as a checklist when wiring your training script and logging.
- **Attention & Tiny Transformers Handout** [5] — detailed derivations and implementation roadmap for the tiny transformer LM that anchors the capstone.

- **Program Schedule** (`schedule.html`) [5] — Week 3 module overview, deadlines, and how this capstone fits into the broader TAE journey.

## 3 What You Will Learn

The bullets below summarize the core skills and concepts to prioritize during Week 3.

- **Attention mechanics** — content-based addressing, scaled dot-product attention, masking, and the transition from single-head to multi-head attention.
- **Transformer blocks** — self-attention, position-wise feedforward networks, residual connections, and layer normalization.
- **Tiny decoder-only LMs** — token and positional embeddings, stacked transformer blocks, and a projection head for next-token prediction.
- **Training and evaluation** — batching, loss shaping for cross-entropy, tracking training/-validation loss, and basic qualitative sampling.
- **LLM engineering habits** — configuration via code or CLI flags, seeding, logging, checkpoints, and simple run records inspired by Part II of the Engineering book.

## 4 How to Study

### Full-Time Track (3–4 h/day)

If you have several hours per day available, use this structure as a default and adapt it to your context.

- **Morning (90 min):** LLM book sections on the overall pipeline and transformer architecture [1]; skim the corresponding slides while keeping the big picture in mind.
- **Midday (60–90 min):** Attention & Tiny Transformers handout [6]; work through scaled dot-product attention, self-attention, and transformer blocks, checking shapes carefully.
- **Afternoon (60 min):** Capstone notebook work — implement or refine `scaled_dot_product_attention`, a transformer block, and the tiny LM forward pass.
- **End (15–30 min):** Engineering hygiene (Part II of the Engineering book [2]) — add logging, configuration, and a short run record (what you changed, what happened).

### Time-Constrained (60–90 min/day)

If your time is limited, use the following lightweight progression to keep moving the capstone forward.

- **Day 1:** Skim LLM overview (book + slides) for global structure; read the handout’s motivation and goal and “How to Use This Capstone” [1, 6].
- **Day 2:** Study scaled dot-product attention and implement `scaled_dot_product_attention` in Colab with the numeric example from the handout [6].
- **Day 3:** Implement single-head self-attention and a minimal transformer block (attention + feedforward + residual + norm) [6].
- **Day 4:** Assemble the tiny decoder-only LM, wire up the loss and training loop, and run a short training session [6, 1].
- **Day 5:** Add simple logging and checkpoints inspired by Part II of the Engineering book; run sampling experiments at different temperatures and document the results [2, 6].

## 5 Focus First (Priorities)

When in doubt, return to these priorities before adding additional features or experiments.

- **Shapes and masking:** whenever you get lost, write down tensor shapes for  $Q, K, V$ , the attention scores, and outputs; verify causal masks by hand on tiny examples [6].
- **Single forward path:** get one clean forward pass from token indices to logits working before adding bells and whistles [1, 6].
- **Training loop basics:** stick to a simple, well-tested pattern (batch → forward → loss → backward → step); log training and validation loss periodically [1, 2].
- **Engineering hygiene:** keep configuration in one place, seed randomness, and save at least one checkpoint per working run; a small run log is more valuable than an extra experiment [2].
- **Capstone alignment:** periodically cross-check your notebook against the handout’s “Colab Implementation Roadmap” and capstone description so you do not drift from the intended scope [6].

## 6 Key Concepts — Concise Recap

### Attention and Self-Attention

Use these bullets as a quick recall checklist while implementing or debugging attention.

- Attention as content-based addressing: queries, keys, values, softmax weights, and weighted sums [6].
- Scaled dot-product attention with masking: score matrix, softmax over keys, and causal masks that prevent looking into the future [6].
- Self-attention:  $Q, K, V$  derived from the same sequence of token embeddings, applied in parallel across positions [6].

### Transformer Blocks

These points summarize what each transformer block is responsible for in your tiny LM.

- Multi-head attention: split heads, compute attention per head, then concatenate and project back to the model dimension [6].
- Position-wise feedforward nets: simple MLP applied independently at each position, sandwiched between residual connections and layer normalization [6].
- Pre-LN vs. Post-LN: for this capstone, a pre-layer normalization structure improves stability at small scales [6].

### Tiny Decoder-Only LM

Keep the following structure in mind when you reason about the end-to-end model.

- Token and positional embeddings combined into a sequence of hidden states of shape  $(B, T, d_{\text{model}})$  [6, 1].
- A stack of transformer blocks followed by a projection head to vocabulary logits; reshape logits and targets for cross-entropy [6, 1].

- Autoregressive sampling: repeatedly feed back generated tokens, using temperature or top- $k$  sampling to control diversity [6, 1].

## Engineering and Tooling

Ground your capstone in these minimal but important engineering practices.

- Configuration objects or CLI flags for hyperparameters; keep them in version control alongside code [2].
- Lightweight tracking (CSV/JSON logs, simple run IDs) to compare experiments; avoid opaque, one-off notebook states [2].
- Checkpoints via `state_dict()` and clear naming conventions so you can resume or re-sample later [2, 1].

## 7 Review Questions

Answer in code or 1–3 sentences; use them as prompts for tiny experiments.

- Write down the shapes of  $Q, K, V$ , the attention score matrix, and the attention output for a batch of size  $B$  and sequence length  $T$ . Where does the scaling factor  $1/\sqrt{d_k}$  appear?
- Explain why causal masking is required for autoregressive language models. How would you test that your mask is implemented correctly?
- In your tiny transformer, where do residual connections appear and why do we use layer normalization around them?
- Show how you reshape logits and targets before applying `nn.CrossEntropyLoss`. What happens if you forget to reshape?
- Describe one engineering decision you made inspired by Part II of the Engineering book (e.g., how you log metrics, how you structure configuration, or how you name checkpoints).
- After training, compare samples at two different temperatures. What changes in the generated text and why?

## 8 Practice Checklist

Use this checklist to confirm that you have covered the essential implementation and engineering pieces for the Week 3 capstone.

- Implement `scaled_dot_product_attention` and test it against the numeric example in the handout [6].
- Build a minimal transformer block with single-head self-attention and a position-wise feedforward network; verify shape invariants end-to-end [6].
- Assemble a tiny decoder-only LM (`TinyTransformerLM`) with embeddings, a small stack of blocks, and a projection head; run a short training loop on a tiny corpus [6, 1].
- Add configuration, logging, and checkpoints following the Engineering book (Part II) and verify that you can reproduce at least one run [2].
- Implement greedy and temperature-based sampling; save a small “gallery” of prompts and completions as part of your capstone artifact [6, 1].

## 9 Tiny Study Loop




Figure 1: Short, repeatable loop for Week 3: read a small piece, run it, inspect shapes/loss/samples, repeat.

## 10 Daily Cadence (Example)

Block	Focus	Resource
45–60 min	Attention/transformer theory	Handout ( <a href="#">att</a> )
20–30 min	LLM pipeline context	<a href="#">llm.html</a> , <a href="#">llm_slides</a>
10–15 min	Engineering hygiene	<a href="#">eng.html</a> (Part II)

## 11 Capstone Project (Week 3)

Build a **Colab-ready tiny transformer language model** that follows the Attention & Tiny Transformers handout and the LLM book patterns. Your model should implement scaled dot-product attention, self-attention, transformer blocks, and a decoder-only architecture that performs next-token prediction on a small text corpus. Host your notebook and any supporting code on **GitHub**; keep the model intentionally small but complete, with seeds, checkpoints, and simple sampling utilities.

- **Execution:** one notebook (plus optional small module) that runs top-to-bottom on Colab; deterministic seeds; clear configuration and device handling.
- **Scope:** tiny decoder-only transformer LM on a simple character- or token-level dataset, with at least one training run that produces non-trivial samples; single or few transformer blocks are sufficient.
- **Deliverables:** GitHub link to the notebook/code; a short run log (what configuration you used, how training behaved, and example samples); optional exported PDF of the handout-aligned results.
- **Schedule:** Reserve **3 × 90 min** focused sessions to (1) build the attention and block components, (2) assemble and train the tiny LM, and (3) refine engineering hygiene and sampling, aligning with the handout’s implementation roadmap.

## References

- [1] Building a Large Language Model from Scratch. [llm.html](#)
- [2] AI, ML & Software Engineering — Building Intelligent Systems in Practice. [eng.html](#)
- [3] Building a Large Language Model from Scratch — Slide Deck. [llm\\_slides.pdf](#)

[4] AI, ML & Software Engineering — Slide Deck (Part II). [eng\\_slides.pdf](#)


[5] Program Schedule. [schedule.html](#)

[6] Attention & Tiny Transformers Handout (From Scaled Dot-Product Attention to a Mini LLM). [att\\_transformers.pdf](#)

# Contact

TAE Week 3 Coaching Guide

LLMs, Attention & Tiny Transformers Capstone



Get in touch:

<https://linktr.ee/dyjh>

© 2025 Dr. Yves J. Hilpisch — All rights reserved.